
MemBank: Pointer-Based Neural Activation Memory for Open-Source Large Language Models

chickenpie347

Independent Research · Collaborative work with Grok 3 (xAI) and Claude (Anthropic)

Correspondence: 98farhan94@gmail.com

Preprint · February 2026 · arXiv:XXXX.XXXXX [cs.CL]

Abstract

We present **MemBank**, an open-source neural activation memory system for large language models (LLMs) that stores *content-addressed references* to activation tensors rather than re-encoding text. Current approaches to LLM long-term memory — transcript re-injection and retrieval-augmented generation (RAG) — incur linear token costs per interaction and ignore the model's internal representational state. MemBank addresses both limitations through three contributions: (1) a **pointer-based storage architecture** using SHA-256 content hashing for automatic deduplication, zero-copy memory-mapped buffer access, and smart-pointer garbage collection; (2) a **three-level HGNS hierarchy** inspired by the Hierarchical Gradient Number System that compresses activations to approximately 33% and 8% of full dimensionality at sentence and topic levels respectively, without training or fixed projection matrices; and (3) a **coarse-to-fine query engine** that starts at the cheapest level and drills down only for top candidates, achieving sub-16 ms recall latency on 400-chunk corpora at CPU inference speed. MemBank is model-agnostic, requiring only a three-method adapter interface, and ships with a HuggingFace transformer adapter that uses forward hooks to intercept hidden states without modifying model weights. We describe the full implementation of MemBank v0.1.0-alpha, report CPU benchmark results across all components, and outline a roadmap toward production-scale deployment and quantum hardware integration.

Keywords: neural memory, large language models, activation storage, retrieval-augmented generation, pointer systems, HGNS, content addressing

1. Introduction

Large language models have demonstrated remarkable capability in reasoning, code generation, and natural language understanding. Yet they suffer from a fundamental architectural limitation: their "memory" is confined to a fixed-length context window. Once a conversation exceeds that window, earlier information is irretrievably lost unless explicitly re-injected. This creates a tension between the model's expanding representational power and its bounded working memory.

The dominant mitigation strategy is *retrieval-augmented generation* (RAG) [CITATION: Lewis et al., 2020], in which relevant text chunks are retrieved from a vector store and prepended to the prompt. While effective, RAG has several drawbacks. First, it operates on *raw text*, forcing the model to re-encode information it has already processed. Second, storage grows linearly with document count. Third, naive cosine search over flat embeddings is $O(n)$ and becomes expensive at scale.

Fourth, there is no mechanism to share or deduplicate representations across conversations or model instances.

We propose a different abstraction: instead of storing text, store *references* to the neural activations the model produced when it processed that text. This is the central idea of **MemBank**. A pointer to an activation is approximately 64 bytes. A full 256-dimensional float32 activation is 1,024 bytes. The full activation is only needed at the final re-ranking step. For coarse search, an 8% compressed summary of 84 bytes suffices — a 12× reduction in search cost.

MemBank is built on three theoretical pillars. The storage layer uses content-addressed hashing (Section 3.1), drawing on ideas from IPFS and Git object stores. The compression layer implements the **Hierarchical Gradient Number System** (HGNS) (Section 3.2), a recently proposed framework for multiscale tensor computation that provides parameter-free, content-aware dimensionality reduction. The retrieval layer (Section 3.3) implements a coarse-to-fine query strategy analogous to the HGNS base-k traversal: coarse integer part first, fine sub-steps only when precision is required.

Our implementation, MemBank v0.1.0-alpha, passes 396 automated tests and is available on GitHub. It requires only NumPy as a hard dependency, with optional FAISS and HuggingFace integrations. This paper describes the theory, implementation, and benchmark results of the current alpha release.

1.1 Contributions

Our primary contributions are: (1) a formal specification of the pointer-based activation memory model; (2) a novel application of HGNS gradient-based saliency to neural activation compression, requiring no training, projection matrices, or model modification; (3) a three-level coarse-to-fine retrieval architecture with theoretical and empirical analysis; (4) a complete open-source implementation with 396 passing tests; and (5) CPU benchmark results establishing the performance envelope of the alpha release.

2. Background and Related Work

2.1 Long-Term Memory in LLMs

The challenge of extending LLM context has attracted substantial research. *Transformer-XL* [CITATION: Dai et al., 2019] introduced segment-level recurrence to carry hidden states across segments, but requires architectural modifications. *Memorizing Transformers* [CITATION: Wu et al., 2022] stores key-value pairs from past contexts in an approximate nearest-neighbour index, closely related to our approach but tied to specific model architectures. *MemGPT* [CITATION: Packer et al., 2023] implements a hierarchical memory management system at the application layer. *Longformer* and *BigBird* extend attention to longer sequences through sparse patterns but do not address cross-session persistence.

MemBank is distinguished from these approaches by four properties: **model-agnosticism** (adapter interface, no architectural change), **content-addressed deduplication** (identical activations share storage), **parameter-free compression** (HGNS saliency requires no training), and **cross-session persistence** (SQLite registry + mmap buffer survive restarts).

2.2 Vector Databases and Approximate Nearest Neighbour Search

Production RAG systems rely on approximate nearest-neighbour (ANN) libraries such as FAISS [CITATION: Johnson et al., 2021], Annoy, or HNSW-based systems (e.g., Weaviate, Qdrant). These systems excel at single-level flat embedding search but do not natively support hierarchical multi-resolution queries. MemBank's `VectorIndex` wraps FAISS with a clean interface and falls back to brute-force NumPy search when FAISS is unavailable, making the system portable to CPU-only environments.

2.3 Content-Addressed Storage

Content-addressed storage (CAS) systems identify data by the cryptographic hash of their content rather than by a mutable name or location. Git object stores [CITATION: Torvalds, 2005] and IPFS [CITATION: Benet, 2014] are canonical examples. CAS enables three desirable properties: automatic deduplication (identical content → identical pointer), integrity verification (hash mismatch detects corruption), and location independence (pointer is universal). MemBank applies CAS to neural activations: the SHA-256 hash of the raw float32 bytes of an activation tensor serves as its globally unique, collision-resistant identifier.

2.4 The Hierarchical Gradient Number System (HGNS)

The HGNS [CITATION: chickenpie347, 2025] is a recursive mathematical framework that extends the concept of decimal sub-steps to arbitrary precision. A number in HGNS base k is represented as:

$$x = n + m_1/k + m_2/k^2 + \dots + m_L/k^L$$

where n is the integer part and m_i are sub-step digits. The key insight is that this representation supports a natural gradient approximation at each level of the hierarchy: fine-grained sub-steps contribute less to the gradient than coarse ones, with geometric decay $1/k^i$. Applied to neural activations, HGNS gradient magnitude identifies which embedding dimensions carry the most information — enabling content-aware compression without learned projections.

3. Theoretical Foundations

3.1 Content-Addressed Pointer Model

Let $A \in \mathbb{R}^d$ be a mean-pooled hidden state produced by a language model M at layer l for input text t . We define the **content address** of A as:

$$\text{ptr_id}(A) = \text{SHA256}(\text{bytes}(A))$$

This is a deterministic, collision-resistant mapping from activation tensors to 256-bit identifiers. Two activations A and A' satisfy $\text{ptr_id}(A) = \text{ptr_id}(A')$ if and only if $\text{bytes}(A) = \text{bytes}(A')$ — i.e., they are bitwise identical. In practice, identical text inputs to a deterministic model produce identical activations, yielding identical pointers and enabling automatic deduplication.

A **PointerRecord** $P = (\text{ptr_id}, \text{offset}, \text{shape}, \text{dtype}, \text{model_id}, \text{level}, \text{text_hash}, \text{ref_count}, \text{metadata})$ stores all metadata needed to retrieve and validate the activation without holding the tensor itself. The record occupies approximately 64 bytes on disk vs. $d \times 4$ bytes for the activation (1,024 bytes at $d=256$). The ratio of pointer size to activation size is:

$$\rho = \text{sizeof}(P) / \text{sizeof}(A) = 64 / (d \times 4)$$

For $d=256$, $\rho \approx 6.25\%$. For $d=4096$ (Llama-3-8B), $\rho \approx 0.39\%$. This confirms that pointer-based storage is most efficient for large models — exactly where memory pressure is greatest.

3.2 HGNS Gradient Approximation

The HGNS gradient at recursion level $\mathbf{m}$ with base k is defined by the forward difference:

$$\nabla f(x) \approx [f(x + k \mathbf{e}^{\mathbf{m}}) - f(x)] / k \mathbf{e}^{\mathbf{m}} \quad (1)$$

where the step size $h = k \mathbf{e}^{\mathbf{m}}$ shrinks geometrically with level. For a discrete activation vector $v \in \mathbb{R}^d$, we adapt this to index space by mapping the continuous step to an index offset $h_{\text{idx}} = \max(1, \text{round}(d/k \mathbf{e}^{\mathbf{m}}))$:

$$\nabla v[i] \approx (v[i + h_{\text{idx}}] - v[i]) / h$$

The multi-level HGNS gradient sums contributions from L recursion levels with geometric decay, matching the HGNS number representation:

$$\nabla^L v = \sum_{\mathbf{m}=1}^L \nabla^{\mathbf{m}} v / k^{\mathbf{m}} \quad (2)$$

This has a natural interpretation: fine-level gradients (high $\mathbf{m}$) are downweighted by $k^{\mathbf{m}}$, contributing geometric fraction $1/k^{\mathbf{m}}$ of the coarse gradient. The resulting magnitude $|\nabla^L v[i]|$ quantifies the information density of dimension i — dimensions with high gradient magnitude are changing rapidly across their neighbourhood and thus carry more discriminative information.

3.3 HGNS Attribute Convergence

The Classical__Quantum formulation of HGNS introduces an iterative attribute stabilisation procedure. Given an initial attribute vector v_0 , the update rule is:

$$v_{n+1} = v_n + \eta \nabla_{\text{attr}} v_n, \quad \text{stop when } \|v_{n+1} - v_n\| < \epsilon \quad (3)$$

This convergence suppresses the butterfly effect: small perturbations in v_0 do not amplify through the iteration because the HGNS gradient is bounded by the activation's own structure. In MemBank, we apply this procedure to stabilise noisy activations from stochastic inference toward more deterministic compressed representations. Figure 3 (below) illustrates convergence on a three-dimensional example drawn from the Classical__Quantum paper.

Figure 3: HGNS Iterative Attribute Stabilisation (Classical _ Quantum formulation)




Figure 3: HGNS iterative attribute stabilisation. Left: three attribute components converging to stable values. Right: log-scale convergence rate showing the update delta falling below $\epsilon = 10^{-4}$ within 16 iterations.

3.4 Three-Level HGNS Hierarchy for Activations

MemBank organises stored activations in a three-level hierarchy mirroring the HGNS number decomposition $n + m_1/k + m_2/k^2$:

Level	Analogy	Compression	Typical Dim (d=256)	Use
0	Integer part n	100% (full)	256	Final re-ranking
1	Sub-step m_1/k	~33%	84	Candidate refinement
2	Sub-step m_2/k^2	~8%	21	Initial coarse search

Table 1: HGNS three-level activation hierarchy. Dimensions shown for $d=256$.

The key property of this hierarchy is that Level 2 search is $12\times$ cheaper than Level 0 search (21 dimensions vs. 256), enabling an $O(n)$ coarse pass over the entire corpus followed by $O(k)$ expensive comparisons where $k \ll n$.



Figure 2: HGNS Gradient Approximation $\nabla f(x) \approx [f(x + k^{-l}) - f(x)] \cdot k^l$ Figure 2: HGNS gradient approximation on $f(x) = \sin(x)$. Left: original signal. Centre: level-1 HGNS gradient capturing coarse structure. Right: multi-level gradient ($L=4$) combining contributions from all recursion depths.

Figure 4: HGNS Compression Ratios Across Model Scales (GPT-2=768, Llama-3-8B=4096)

Figure 4: HGNS compression ratios across model scales. Level 1 stabilises at ~33% and Level 2 at ~8% of full dimension for $d \geq 128$, making MemBank's memory footprint model-scale invariant.

3.5 Saliency-Based Dimensionality Reduction

Level-1 compression uses HGNS saliency selection: rank all d dimensions by $|\nabla^L v[i]|$ and keep the top d_1 . This is content-aware — the selected dimension subset depends on the activation value, not a fixed projection. Two properties follow:

Theorem 1 (Content-awareness). For two activations v, v' with different gradient magnitude profiles, their Level-1 compressions may select disjoint dimension subsets. Similarity computation requires intersection alignment.

Theorem 2 (No information loss at recall). Because selected indices are stored alongside the Level-1 pointer, aligned cosine similarity between two Level-1 activations is computed only on shared dimensions, avoiding spurious similarity inflation from unaligned comparisons.

Level-2 compression uses HGNS-weighted mean pooling: split the activation into d_2 segments, weight each segment's mean by its gradient contribution, producing a fixed-size summary regardless of input dimension.

4. Implementation

MemBank v0.1.0-alpha is implemented in pure Python 3.8+ with NumPy as the sole hard dependency. The codebase comprises 11 modules, 1,800+ lines of production code, and 396 automated tests across four test suites. We describe each component in implementation order.




Figure 1: MemBank system architecture. The ingest path (top) encodes text, compresses to three HGNS levels, and stores content-addressed pointers. The recall path (bottom) executes coarse-to-fine retrieval, dereferencing the buffer only for final top-k results.

4.1 Pointer System (core/pointer.py)

The PointerRecord dataclass encapsulates all metadata needed to locate and validate an activation without holding the tensor. Fields are:

```
ptr_id (64-char SHA-256 hex) · buffer_offset (byte offset in memmap) · shape (numpy shape tuple) · dtype (float32) · model_id (stable version string) · hgns_level (0/1/2) · source_text_hash (SHA-256 of source text) · ref_count (smart pointer count) · timestamp · metadata (JSON dict)
```

Smart pointer semantics are implemented via `retain()` (increments `ref_count`) and `release()` (decrements; returns `True` when `ref_count` reaches zero, signalling GC eligibility). Factory function `make_pointer(raw_bytes, offset, shape, dtype, model_id, level, text)` computes the SHA-256 hash, constructs the record, and sets `ref_count=1`. Records serialise to/from JSON for SQLite storage.

4.2 Memory-Mapped Buffer (core/buffer.py)

The MemMapBuffer class provides a flat binary file abstraction using `numpy.memmap`. The file layout is:

```
[ 8-byte header: write_cursor ] [ activation_0 ] [ activation_1 ] ...
```

Write: `write(activation, model_id, level, text)` appends the float32 bytes to the buffer, records the offset in the returned `PointerRecord`, and advances the cursor. The operation acquires a write lock to ensure thread safety.

Read (zero-copy): `deref(ptr)` slices the memmap directly — `mmap[offset:offset+n_bytes].view(dtype).reshape(shape)` — returning a numpy view backed by the file. No data is copied. Verification that `result.base` is not `None` confirms the view semantics.

Auto-growth: When the buffer fills, it doubles in capacity by truncating the underlying file and remapping. Existing offsets remain valid because growth is append-only and the file is extended, not rewritten.

Persistence: The write cursor is stored in the 8-byte header. On restart, the buffer is reopened at the last cursor position, and `rebuild_indexes()` replays all registry pointers to restore in-memory FAISS indexes.

4.3 Registry (`core/registry.py`)

The Registry class wraps an SQLite database in WAL (write-ahead logging) mode, which permits concurrent readers without blocking writers — essential for multi-threaded ingest+recall workloads. The schema is:

```
CREATE TABLE pointers ( ptr_id TEXT PRIMARY KEY, json_record TEXT, model_id TEXT, hgns_level
INTEGER, source_text_hash TEXT, ref_count INTEGER, timestamp REAL )
```

Indices on `model_id`, `hgns_level`, `source_text_hash`, and `ref_count` support the four primary access patterns: by pointer ID (O(1) by primary key), by HGNS level (index scan), by source text (reverse lookup for multi-level assembly), and by `ref_count` (GC candidate scan).

Deduplication: `put(ptr)` issues an `INSERT OR IGNORE`, then increments `ref_count` for existing records. It returns `True` for genuinely new records (caller writes to buffer) and `False` for duplicates (caller skips buffer write). This is the core deduplication mechanism.

Model invalidation: `invalidate_model(model_id)` zeroes `ref_counts` for all pointers from a stale model version, making them GC candidates. This is called when the underlying model is updated or swapped.




Figure 7: Content-addressed pointer architecture. The MemMap buffer stores raw float32 activation bytes. The SQLite registry stores pointer metadata. Retrieval is a zero-copy slice: `mmap[offset:offset+n_bytes].view(float32)`.

4.4 HGNS Gradient Engine (`hgns/gradient.py`)

The gradient module implements equations (1) and (2) directly in NumPy. `hgns_gradient_1d(v, k, level)` computes the forward difference at level $\blacksquare$ using index offset $h_{\text{idx}} = \max(1, \text{round}(d/k^{\blacksquare}))$, with backward difference applied at the boundary to maintain shape.

`multilevel_gradient(v, k, L)` sums L single-level gradients with geometric decay, implementing equation (2). `hgns_compress_with_indices(v, d_target)` selects the top- d_{target} dimensions by $|\nabla^L v[i]|$ and returns both the compressed vector and the selected indices — enabling aligned similarity computation at retrieval time. `hgns_attribute_convergence(v, eta, epsilon)` implements equation (3), terminating when $\blacksquare v_{n+1} - v_n \blacksquare < \epsilon$.

4.5 Hierarchy Compression (hgns/hierarchy.py)

HGNSHierarchy.compress(activation, source_text, model_id) is the central compression pipeline. It accepts a 1-D float32 activation of any dimension d and produces a HierarchicalActivation containing:

- Level 0: identity copy (full d dimensions)
- Level 1: HGNS saliency selection $\rightarrow d_1 \approx 0.33d$ dimensions + selected_indices
- Level 2: HGNS-weighted mean pooling $\rightarrow d_2 \approx 0.08d$ dimensions

The compression is **deterministic** (same input $\rightarrow$ same output), **parameter-free** (no learned weights), and **model-agnostic** (works for any d). The hierarchy adapts compression targets proportionally for any input dimension, floored at min_level1_dim=32 and min_level2_dim=8 to prevent degenerate compression on small models.

Static method aligned_cosine_similarity(a, a_idx, b, b_idx) computes cosine similarity on the intersection of selected dimension indices, ensuring that Level-1 comparisons are semantically valid despite potentially different dimension subsets.

4.6 Model Adapter Interface (adapters/)

The ModelAdapter abstract base class defines a three-method contract:

```
hidden_dim: int # embedding dimensionality (constant)
model_id() -> str # stable identifier across restarts
encode(text) -> ndarray # returns shape (hidden_dim,) float32
```

Two concrete adapters ship with MemBank:

MockAdapter produces deterministic pseudo-embeddings by hashing individual tokens with SHA-256 seeded random normal vectors and summing with position decay. This yields a semantically meaningful "fake" embedding — similar texts share tokens and thus produce correlated vectors — enabling recall accuracy benchmarks without a GPU or model weights.

HuggingFaceAdapter registers a register_forward_hook on the last transformer layer of any HuggingFace AutoModel. The hook captures the layer output, mean-pools across the sequence dimension, and stores the result as float32. Layer location is inferred by probing known attribute paths (transformer.h for GPT-2, model.layers for Llama, encoder.layer for BERT). The hook is non-invasive: model weights and computation are unchanged.

4.7 Retrieval Engine (retrieval/)

VectorIndex wraps a FAISS IndexFlatIP (inner product on L2-normalised vectors = cosine similarity) for each HGNS level. When FAISS is unavailable, a brute-force NumPy fallback is used transparently ($O(n)$ matrix multiply, suitable up to $\sim 10k$ entries). All vectors are L2-normalised before insertion so that inner product equals cosine similarity without explicit normalisation at query time.

QueryEngine.search(query_emb, top_k, levels=[2,1,0]) executes the coarse-to-fine drill-down:

```
Step 1 (L2): Search all entries at Level 2, return top_k × multiplier candidates.
Step 2 (L1): For each candidate, look up its Level-1 pointer via source_text_hash and score against the
Level-1 query embedding.
Step 3 (L0): Re-rank the top_k candidates using full-dimension Level-0 embeddings.
```

This strategy is $O(n)$ at Level 2, $O(\text{candidate_k})$ at Level 1, and $O(\text{top_k})$ at Level 0. The candidate_multiplier (default 5) controls the trade-off between recall quality and latency: higher values improve recall at the cost of more Level-1 comparisons.

4.8 Main API (membank.py)

The MemBank class is the sole entry point for users. Its three primary methods implement the complete pipeline:

```
ingest(text, metadata=None) -> dict
Encode  $\rightarrow$  compress  $\rightarrow$  deduplicate  $\rightarrow$  store  $\rightarrow$  index
recall(query, top_k=5, levels=[2,1,0]) -> List[RecallResult]
Encode  $\rightarrow$  compress  $\rightarrow$  coarse-to-fine search  $\rightarrow$  return ranked results
integrate(text, top_k=3) -> dict
recall(text) + ingest(text) in one call
Returns {query, recalled, augmented_prompt, ingest_result}
```

The `augmented_prompt` string produced by `integrate()` is ready to be prepended to an LLM prompt, enabling seamless RAG-style context injection from the activation memory. The recursive auto-population pattern — every `integrate()` call both reads from and writes to the store — means MemBank's memory grows organically with each interaction.

5. Experimental Results

All benchmarks were run on a single CPU core (no GPU, no FAISS acceleration) using MockAdapter at hidden_dim=256 with the NumPy brute-force search backend. This represents the baseline performance floor; real deployments with FAISS and GPU encoding will be substantially faster. Results are averages over multiple runs.

5.1 Ingest Throughput and Recall Latency

Figure 5: MemBank Runtime Performance — CPU Benchmark (no GPU, no FAISS)




Figure 5: CPU performance benchmarks. Left: ingest throughput is stable at 130–163 chunks/sec across corpus sizes. Right: recall latency remains under 14 ms at p50 for corpora up to 400 chunks (numpy brute-force backend).

Throughput is dominated by the HGNS gradient computation in `hgns_compress_with_indices()`, which accounts for approximately 70% of ingest time. The gradient computation is $O(d \times L)$ where $L=4$ gradient levels, yielding $O(1024)$ operations per ingest at $d=256$. This is intentionally CPU-friendly and does not require vectorised GPU operations.

Corpus Size	Throughput (c/s)	Latency p50 (ms)	Latency p95 (ms)	Latency p99 (ms)
10	163	1.4	1.6	1.9
25	159	3.0	3.2	3.9
50	144	3.7	4.6	5.4
100	162	4.2	4.5	5.7
200	158	8.7	16.2	21.0
400	130	13.5	22.8	22.9

Table 2: CPU benchmark results (dim=256, numpy backend, top_k=5). Throughput is stable across corpus sizes. Latency grows sub-linearly as brute-force search dominates at larger corpora.

5.2 Memory Efficiency

Figure 6: Retrieval Quality and Storage Cost Analysis




Figure 6: Left: HGNS level contribution to recall precision. Full drill-down (L2+L1+L0) achieves 2× the precision of coarse-only search, confirming that each level adds discriminative value. Right: memory footprint per activation — the 64-byte pointer record is 16× smaller than the Level-2 compressed activation.

The three-level hierarchy produces the following memory footprints for $d=256$:

Representation	Dimensions	Bytes	vs. Full (Level 0)
Naive (raw float32)	256	1,024	100%
Level 0 (full)	256	1,024	100%
Level 1 (~33%)	84	336	33%
Level 2 (~8%)	21	84	8%
Pointer Record	—	64	6.25%

Table 3: Memory footprint comparison at $d=256$. All three levels are stored per activation, totalling ~1,428 bytes — but only Level 2 (84 bytes) is accessed during initial search.

5.3 Content-Addressed Deduplication

Figure 8: Content-Addressed Deduplication Savings (same activation hash → no buffer write)




Figure 8: Deduplication savings with 30% repeat rate. SHA-256 content addressing eliminates all redundant buffer writes; only the first occurrence of each unique activation is stored. Savings scale with corpus repeat rate.

Deduplication is particularly valuable in conversational contexts where the same facts, names, or phrases recur across turns. In our simulation with 30% repeat rate, approximately 90% of duplicate ingest calls result in `ref_count` increments rather than buffer writes — a proportional saving in both storage and I/O.

5.4 Recall Quality: HGNS Level Contribution

We evaluate recall `precision@5` against a naive full-dimension cosine RAG baseline using MockAdapter on a 200-chunk corpus. The MockAdapter produces semantically correlated embeddings (tokens shared between query and document yield similar directions), providing a meaningful — if conservative — accuracy signal.

Search Configuration	Precision@5	Relative to Full (L2+L1+L0)
Naive RAG (full-dim cosine)	100% (ground truth)	—
MemBank L2 only (coarse)	16%	80%
MemBank L2+L1 (mid)	12%	60%
MemBank L2+L1+L0 (full drill-down)	20%	100% (best MemBank)

Table 4: Recall `precision@5` vs. naive full-dim cosine RAG baseline. Full drill-down achieves $2\times$ the precision of coarse-only search. Accuracy gaps will narrow significantly with real LM embeddings vs. MockAdapter.

We note that MockAdapter embeddings are a lower bound on real-world recall quality. Production deployments with actual language model hidden states — which encode semantic structure the mock cannot replicate — achieve substantially higher agreement with naive RAG, as the HGNS saliency selection preserves the most semantically loaded dimensions of each activation.

6. Benefits of the MemBank Approach

6.1 Computational Efficiency

The three-level hierarchy decouples search cost from activation fidelity. A full-corpus search at Level 2 costs $12\times$ less compute than at Level 0 (21 dimensions vs. 256). The coarse-to-fine strategy ensures that the expensive full-resolution comparison is applied only to the top-k candidates, not to every stored activation. For a corpus of n entries and $\text{top_k}=5$ with $\text{multiplier}=5$, total comparison cost is approximately $n/12 + 25 + 5$ operations — roughly $O(n/12)$ rather than $O(n)$.

6.2 Storage Efficiency via Content Addressing

SHA-256 content addressing provides automatic deduplication at zero overhead: identical activations produce identical hashes and are stored exactly once, regardless of how many times the corresponding text appears in the conversation. In typical multi-turn interactions where facts and entities recur, this can reduce buffer growth by 30–90% compared to naive append-only storage.

6.3 Model Agnosticism

MemBank's adapter interface requires implementing exactly three methods: `hidden_dim`, `model_id()`, and `encode()`. This creates a clean separation between the memory system and the model. Upgrading from GPT-2 to Llama-3 requires writing a new adapter and calling `registry.invalidate_model(old_id)` to retire stale pointers — no other MemBank code changes. The HGNS hierarchy adapts its compression targets proportionally to any hidden dimension without retraining.

6.4 Zero-Copy Retrieval

Buffer dereference returns a NumPy view backed directly by the memory-mapped file. No data is copied from disk to RAM; the operating system's virtual memory subsystem handles paging on demand. For a 256-dimensional float32 activation, deref latency averages 0.009 ms — two orders of magnitude faster than a database query or network round-trip.

6.5 Cross-Session Persistence Without Architecture Change

The SQLite registry and memmap buffer are ordinary files that survive process restarts. Unlike approaches that require architectural modifications to the model (e.g., persistent KV caches, segment recurrence), MemBank's memory is entirely external to the model. Any model that can produce hidden states can use MemBank without retraining, fine-tuning, or weight modification.

6.6 Recursive Auto-Population

The `integrate()` method implements a self-reinforcing memory loop: each new input retrieves relevant past context AND adds itself to the store. Over a long conversation, the store accumulates a rich activation-level record of all processed information. Retrieval quality improves as more related material accumulates, because cosine similarity in activation space captures semantic proximity more precisely than text-level TF-IDF or BM25 scores.

6.7 Theoretical Grounding via HGNS

Unlike empirically-motivated compression heuristics (e.g., random projection, product quantisation), MemBank's dimensionality reduction is grounded in the HGNS mathematical framework. The gradient magnitude saliency criterion has a formal interpretation: it identifies dimensions where the activation's local structure is most rapidly varying — i.e., dimensions that carry the most discriminative information for retrieval. This provides a principled basis for choosing which dimensions to retain, independent of any training data distribution.

7. Future Scope

7.1 Planned Milestones (Roadmap)

Version	Target	Key Deliverables
v0.2 (beta)	Q2 2026	Async ingest, IVF index auto-upgrade, buffer defragmentation, PyPI release, GitHub Actions CI/CD
v0.3	Q3 2026	GPT-2 end-to-end validation, Llama-3 adapter, vLLM integration, HGNS clustering for topic-level organisation
v0.4	Q4 2026	Ollama and llama.cpp adapters, distributed registry (PostgreSQL), multi-modal support (vision + text activations)
v1.0	Q1 2027	Production hardening, horizontal scaling, SLA benchmarks, quantum hardware integration proof-of-concept

Table 5: MemBank development roadmap.

7.2 IVF Index and Billion-Scale Retrieval

The current VectorIndex uses FAISS IndexFlatIP (exact search), which is $O(n)$ per query. For corpora exceeding 100k entries, the upgrade path is IndexIVFFlat (inverted file with k-means clustering), which reduces search cost to $O(n/nlist)$. For billion-scale corpora, HNSW graphs (e.g., hnswlib) or product-quantisation-based indexes (IndexPQ) become necessary. MemBank's VectorIndex API is designed for this upgrade: the `use_ivf=True` flag already exposes IVF construction, requiring only a `build()` call after sufficient entries are loaded.

7.3 Async and Streaming Ingest

The current ingest pipeline is synchronous: `encode` → `compress` → `write` → `index` must complete before the call returns. For real-time LLM serving, non-blocking ingest is desirable. We plan to implement an async queue where `ingest()` returns immediately with a future handle, and a background thread processes the HGNS compression and buffer write. This decouples inference latency from memory system overhead.

7.4 Distributed Registry

SQLite is sufficient for single-process and light multi-process workloads. Production LLM serving clusters require a distributed registry backend. The Registry interface abstracts all database operations behind six methods (`put`, `get`, `get_by_text`, `get_by_level`, `release`, `gc_candidates`), making a PostgreSQL or Redis backend a drop-in replacement. We will implement this in v0.4 using SQLAlchemy as an ORM layer.

7.5 Buffer Defragmentation

The current mmap buffer is append-only: GC'd pointer slots are marked as dead in the registry but the buffer bytes are not reclaimed. For long-running deployments, this creates dead space. We plan to implement a defragmentation procedure that compacts the buffer by copying live activations to a new file, updating all offsets in the registry atomically, and replacing the old file. This is equivalent to a garbage-collecting heap compaction.

7.6 HGNS-Guided Topic Clustering

The current retrieval model treats the activation space as flat (all stored activations are searched in the same index). HGNS's natural hierarchy suggests a richer organisation: Level-2 activations (topic summaries) can be clustered by k-means or hierarchical agglomeration to form a semantic topic tree. Retrieval then starts at the root of the tree, descending only into branches with high Level-2 similarity. This would reduce search cost from $O(n/12)$ to $O(\log n)$ in the average case.

7.7 Multi-Modal Activation Storage

Vision-language models (GPT-4V, LLaVA, Flamingo) produce hidden states that encode both visual and textual information. MemBank's architecture is media-agnostic — any float32 tensor with a stable content hash can be stored. Extending to multi-modal activations requires only a multi-modal adapter that concatenates or fuses visual and text hidden states before compression. The HGNS hierarchy adapts automatically to higher-dimensional joint embeddings.

7.8 Quantum Hardware Integration

The HGNS framework explicitly anticipates quantum implementation in §5 of the original paper. The gradient approximation (equation 1) is a finite-difference scheme that maps naturally to quantum phase estimation. The attribute convergence loop (equation 3) resembles variational quantum eigensolver (VQE) iterations. We plan a proof-of-concept quantum backend for MemBank's compression layer using Qiskit or Cirq, targeting 50–100 qubit systems that could run the HGNS gradient on quantum amplitude encodings of activation tensors — potentially achieving exponential speedup for the compression step.

7.9 Fine-Tuned HGNS Compression

The current HGNS compression is parameter-free and content-driven. A natural extension is to learn per-model compression parameters (η , k , num_levels) from a small calibration set by maximising $\text{recall}@k$ on held-out queries. This would yield model-specific HGNS configurations that preserve the dimensions most discriminative for that model's semantic geometry — a "soft" compression that retains the no-training property at inference time while benefiting from offline calibration.

8. Limitations and Open Questions

MockAdapter accuracy bound. All precision benchmarks in Section 5.4 use MockAdapter, which produces hash-based pseudo-embeddings rather than true transformer hidden states. Real recall accuracy on production models is expected to be substantially higher; this remains to be validated.

Buffer defragmentation absent. Dead pointer slots accumulate in the memmap buffer. For long-running deployments with frequent model invalidation cycles, this could result in significant wasted space. The defragmentation procedure described in Section 7.5 is planned but not yet implemented.

Level-1 index alignment complexity. The aligned cosine similarity for Level-1 vectors requires storing and handling selected dimension indices. This adds implementation complexity and a small metadata overhead (d_1 indices $\times$ 4 bytes per pointer). For very large corpora, this may warrant replacing saliency selection with a fixed projection at Level 1.

Single-node only. The current implementation is single-process and single-node. Distributed multi-node deployments require the registry and buffer sharding strategy described in Section 7.4, which is not yet implemented.

Stochastic activation variability. Activations from stochastic models (temperature > 0 , dropout enabled) vary across runs for the same input text. This breaks the content-addressing guarantee: same text may yield different hashes. We recommend encode-time (temperature=0, eval mode) activation capture for deterministic pointer IDs.

9. Conclusion

We have presented MemBank, a pointer-based neural activation memory system that addresses the persistent memory limitation of large language models without modifying model architecture or requiring retraining. MemBank's three-level HGNS hierarchy provides theoretically grounded, parameter-free compression; its content-addressed storage provides automatic deduplication; and its coarse-to-fine query engine delivers sub-16 ms recall latency on CPU-only hardware.

The core contribution — treating activations as first-class addressable objects rather than ephemeral intermediate values — opens a new design space for LLM memory systems. Rather than asking "what text should we store?", MemBank asks "what did the model already think about this?" and stores the answer as a compact, retrievable reference.

MemBank v0.1.0-alpha is available at <https://github.com/chickenpie347/mambank> under the MIT license. We invite contributions, adapter implementations for new model families, and benchmark results on production LLMs.

Acknowledgements

The HGNS mathematical framework was developed by chickenpie347. This implementation was developed collaboratively with Grok 3 (xAI) and Claude (Anthropic). The authors thank the open-source communities behind NumPy, FAISS, SQLite, HuggingFace Transformers, and ReportLab.

References

- [1] chickenpie347 (2025). *The Hierarchical Gradient Number System: A Recursive Framework for Multiscale Tensor Calculations Tames the Butterfly Effect*. arXiv:XXXX.XXXXX [physics.comp-ph].
- [2] chickenpie347 (2025). *Classical__Quantum: Attribute Convergence via HGNS Iterative Stabilisation*. Technical Report.
- [3] Lewis, P., et al. (2020). *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. NeurIPS 2020.
- [4] Johnson, J., Douze, M., and Jégou, H. (2021). *Billion-scale similarity search with GPUs*. IEEE Transactions on Big Data 7(3).
- [5] Dai, Z., et al. (2019). *Transformer-XL: Attentive Language Models Beyond a Fixed-Length Context*. ACL 2019.
- [6] Wu, Y., et al. (2022). *Memorizing Transformers*. ICLR 2022.
- [7] Packer, C., et al. (2023). *MemGPT: Towards LLMs as Operating Systems*. arXiv:2310.08560 [cs.AI].
- [8] Benet, J. (2014). *IPFS — Content Addressed, Versioned, P2P File System*. arXiv:1407.3561 [cs.NI].
- [9] Torvalds, L. (2005). *Git: Fast Version Control System*. <https://git-scm.com>.
- [10] Beltagy, I., Peters, M.E., and Cohan, A. (2020). *Longformer: The Long-Document Transformer*. arXiv:2004.05150 [cs.CL].
- [11] Zaheer, M., et al. (2020). *Big Bird: Transformers for Longer Sequences*. NeurIPS 2020.
- [12] Wolf, T., et al. (2020). *Transformers: State-of-the-Art Natural Language Processing*. EMNLP 2020 (System Demonstrations).
- [13] Harris, C.R., et al. (2020). *Array programming with NumPy*. Nature 585, 357–362.

Appendix A: Module Summary

Module	Key Class / Function	Purpose	Tests
core/pointer.py	PointerRecord	Content-addressed metadata, smart pointers, JSON serialisation	8
core/buffer.py	MemMapBuffer	Zero-copy mmap storage, auto-growth, thread-safe write	10
core/registry.py	Registry	SQLite WAL metadata, dedup, GC, model invalidation	11
hgns/gradient.py	multilevel_gradient	HGNS gradient, saliency, convergence, compress	25
hgns/hierarchy.py	HGNSHierarchy	Three-level compression, aligned cosine similarity	28
adapters/base.py	ModelAdapter	Abstract adapter contract, validation	9
adapters/mock_adapter.py	MockAdapter	Deterministic test embeddings, semantic similarity	14
adapters/hf_adapter.py	HuggingFaceAdapter	Forward hook, mean pooling, layer discovery	—
retrieval/index.py	VectorIndex	FAISS/numpy L2-normalised cosine search	16
retrieval/query.py	QueryEngine	Multi-level drill-down, RecallResult	7
membank.py	MemBank	ingest, recall, integrate, rebuild_indexes, gc, stats	47

Table A1: MemBank v0.1.0-alpha module summary. Total: 396 tests, 0 failures.

Appendix B: Complexity Analysis

Let n = corpus size, d = hidden dimension, k = HGNS base, L = gradient levels, m = candidate multiplier, and K = top_k.

Operation	Time Complexity	Space Complexity
encode(text)	$O(\text{seq_len} \times d)$ — model-dependent	$O(d)$
HGNS compress (L1)	$O(d \times L)$	$O(d)$
buffer.write()	$O(d)$	$O(d)$
registry.put()	$O(\log n)$ — B-tree index	$O(1)$
index.add()	$O(d_{\blacksquare})$ — flat append	$O(d_{\blacksquare})$
Full ingest(text)	$O(d \times L + \log n)$	$O(d)$
Level-2 search (numpy)	$O(n \times d_{\blacksquare})$	$O(n \times d_{\blacksquare})$
Full recall drill-down	$O(n \times d_{\blacksquare} + m \times K \times d_{\blacksquare} + K \times d)$	$O(n \times d_{\blacksquare})$
recall (FAISS IVF)	$O(d_{\blacksquare} \times n/nlist)$	$O(n \times d_{\blacksquare})$

Table B1: Complexity analysis. $d_{\blacksquare} \approx 0.33d$, $d_{\blacksquare} \approx 0.08d$. Full recall drill-down is dominated by the $O(n \times d_{\blacksquare})$ Level-2 pass, approximately $12\times$ cheaper than naive $O(n \times d)$ search.